


[learn-co-curriculum](#) / [dsc-pivot-tables-pandas-lab](#) Public View license 0 stars  201 forks Star Watch ▾[Code](#) [Issues](#) [Pull requests](#) 2 [Actions](#) [Projects](#) [Security](#) [Insights](#) solution ▾

...

This branch is [8 commits ahead](#), [8 commits behind](#) master. Contribute ▾sumedh10 Merge pull request [#2](#) from learn-co-curriculum/v2-1 ...

on Oct 17, 2019

 10[View code](#) README.md

Pivot Tables with Pandas - Lab

Introduction

In this lab, we'll learn how to make use of our newfound knowledge of pivot tables to work with real-world data.

Objectives

In this lab you will:

- Describe what is meant by long and wide format data
- Use multi-hierarchical indexing to access aggregated data
- Use pivot to create a more organized aggregated DataFrame
- Use stack and unstack to move between different level of multi-indexing

Getting Started

In the cell below:

- Import `pandas` and set the standard alias
- Import `matplotlib.pyplot` and set the standard alias
- Run the `iPython` magic command to display `matplotlib` graphs inline within the notebook

```
import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
```

Load the data

The data for this activity is stored in a file called `'causes_of_death.tsv'` which is a somewhat morbid dataset from the center for disease control. Note that the file extension `.tsv` indicates that this data is formatted slightly differently than the standard `.csv`, the difference being that it has 'tab separated values' instead of 'comma separated values'. As such, pass in the optional parameter `delimiter='\t'` into the `pd.read_csv()` function.

```
df = pd.read_csv('causes_of_death.tsv', delimiter='\t')
```

Now, display the head of the `DataFrame` to ensure everything loaded correctly.

```
df.head()
```

<style scoped> .dataframe tbody tr th:only-of-type { vertical-align: middle; }

```
.dataframe tbody tr th {
    vertical-align: top;
}
```

```
.dataframe thead th {
    text-align: right;
}
```

</style>

	Notes	State	State Code	Ten-Year Age Groups	Ten-Year Age Groups Code	Gender	Gender Code	Race
0	NaN	Alabama	1	< 1 year	1	Female	F	American Indian or Alaska Native
1	NaN	Alabama	1	< 1 year	1	Female	F	Asian or Pacific Islander
2	NaN	Alabama	1	< 1 year	1	Female	F	Black or African American
3	NaN	Alabama	1	< 1 year	1	Female	F	White
4	NaN	Alabama	1	< 1 year	1	Male	M	Asian or Pacific Islander

Our data is currently in **Wide** format. We can tidy this up by converting it to **Long** format by using groupby statements to aggregate our data into a much neater, more readable format.

Groupby aggregations

Complete the following groupby statements.

- Groupby `State` and `Gender` . Sum the values.

```
# Your code here
df.groupby(['State', 'Gender']).sum().head()
```

<style scoped> .dataframe tbody tr th:only-of-type { vertical-align: middle; }

```
.dataframe tbody tr th {
  vertical-align: top;
}

.dataframe thead th {
  text-align: right;
}
```

</style>

		Notes	State Code	Deaths
State	Gender			
Alabama	Female	0.0	40	430133
	Male	0.0	41	430647
Alaska	Female	0.0	80	27199
	Male	0.0	84	36135
Arizona	Female	0.0	180	396028

- Groupby State , Gender , and Race . Find the average values.

```
# Your code here
df.groupby(['State', 'Gender', 'Race']).mean().head()
```

<style scoped> .dataframe tbody tr th:only-of-type { vertical-align: middle; }

```
.dataframe tbody tr th {
  vertical-align: top;
}

.dataframe thead th {
  text-align: right;
}
```

</style>

			Notes	State Code	Deaths
State	Gender	Race			

			Notes	State Code	Deaths
State	Gender	Race			
Alabama	Female	American Indian or Alaska Native	NaN	1.0	70.875000
		Asian or Pacific Islander	NaN	1.0	95.500000
		Black or African American	NaN	1.0	9074.000000
		White	NaN	1.0	29890.636364
	Male	American Indian or Alaska Native	NaN	1.0	86.375000

- Groupby Gender and Race . Find the minimum values.

```
# Your code here
df.groupby(['Gender', 'Race']).min().head()
```

```
<style scoped> .dataframe tbody tr th:only-of-type { vertical-align: middle; }
```

```
.dataframe tbody tr th {
    vertical-align: top;
}
```

```
.dataframe thead th {
    text-align: right;
}
```

```
</style>
```

		Notes	State	State Code	Ten-Year Age Groups	Ten-Year Age Groups Code	Gender Code
Gender	Race						

		Notes	State	State Code	Ten-Year Age Groups	Ten-Year Age Groups Code	Gender Code
Gender	Race						
Female	American Indian or Alaska Native	NaN	Alabama	1	1-4 years	1	F
	Asian or Pacific Islander	NaN	Alabama	1	1-4 years	1	F
	Black or African American	NaN	Alabama	1	1-4 years	1	F
	White	NaN	Alabama	1	1-4 years	1	F
Male	American Indian or Alaska Native	NaN	Alabama	1	1-4 years	1	M

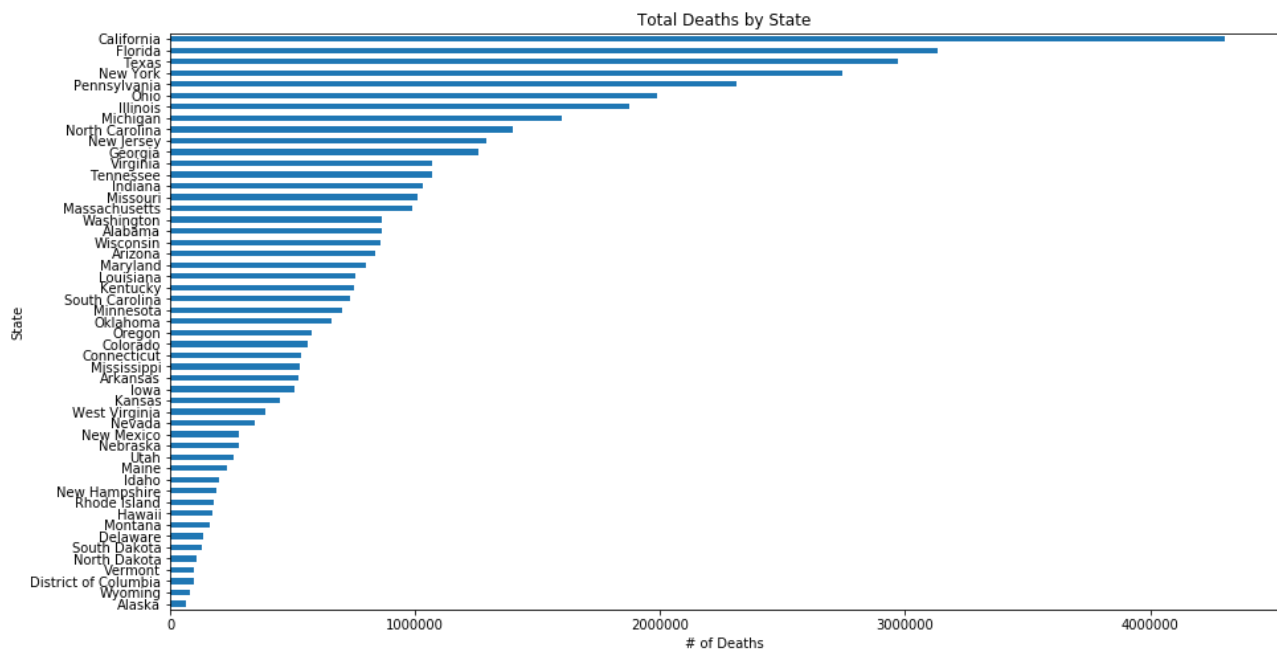
Create a bar chart of the total number of deaths by state:

- Sort your columns in order (ascending or descending are both acceptable).
- Also make sure to include a title, axes labels and have your graph be an appropriate size.

NOTE: In order to do this, slice the `Deaths` column after the `.groupby()` method, but before the `.sum()` method. You can even chain the `.plot()` method on after the `.sum()` method and do this all on one line, excluding the labeling of the graph!

```
# Your code here
df.groupby(['State'])['Deaths'].sum().sort_values().plot(kind='barh', figsize=(15,8))
plt.title('Total Deaths by State')
plt.xlabel("# of Deaths")
```

```
Text(0.5, 0, '# of Deaths')
```



Inspecting our data

Let's go one step further and print the data type of each column.

In the cell below, use the `.info()` method of the DataFrame, and note the data type that each column is currently stored as.

```
df.info()
```

Let's look at some samples from the `Population` column to see if the current encoding seems appropriate for the data it contains.

In the cell below, display the population values for the first 5 rows in the DataFrame.

```
df.Population.iloc[:5]
```

```
0    3579
1    7443
2   169339
3   347921
```

```
4      7366
```

```
Name: Population, dtype: object
```

Just to be extra sure, let's check the value counts to see how many times each unique value shows up in the dataset. We'll only look at the top 5.

In the cell below, print out the top 5 `value_counts()` of the population column of the DataFrame.

```
df.Population.value_counts()[:5]
```

```
Not Applicable    75
```

```
14810              2
```

```
113598             2
```

```
11680              2
```

```
6420               2
```

```
Name: Population, dtype: int64
```

Clearly, this data should be stored as a numeric type, not a categorical type.

Reformat the Population column as an integer

As it stands, not all values can be reformatted as integers. Most of the cells in the `Population` column contain integer values, but the entire column is currently encoded in string format because some cells contain the string `'Not Applicable'`.

We need to remove these rows before we can cast the `Population` column to an integer data type.

In the cell below:

- Slice the rows of `df` where the `Population` column is equal to `'Not Applicable'`
- Use `to_drop.index` to drop the offending rows from `df`. Be sure to set the `axis=0`, and `inplace=True`
- Cast the `Population` column to an integer data type using the `.astype()` method, with the single parameter `int64` passed in
- Print the `Population` column's `dtype` attribute to confirm it is now stored in `int64` format

NOTE: `.astype()` returns a copy of the column, so make sure you set the `Population` column equal to what this method returns--don't just call it!


```
# Your code here
to_drop = df[df['Population'] == 'Not Applicable']
df.drop(to_drop.index, axis=0, inplace=True)
df['Population'] = df['Population'].astype('int64')
print(df['Population'].dtype)
```


```
int64
```

Complete the bar chart

Now that we've reformatted our data, let's create a bar chart of the mean Population by State .

```
# Your code here
df.groupby('State')['Population'].mean().plot(kind='barh', figsize=(15,8))
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x120e7fa90>
```



Below we will investigate how we can combine the `.pivot()` method along with the `.groupby()` method to combine some cool **stacked bar charts**!

Use aggregate methods

In the cell below:

- Group `df` by `'State'` and `'Gender'` , and then slice both `'Deaths'` and `'Population'` from it. Chain the `.agg()` method to return the mean, min, max, and standard deviation of these sliced columns.

NOTE: This only requires one line of code.

By now, you've probably caught on that the code required to do this follows this pattern:

```
[things to group by][columns to slice].agg([aggregates to return])
```

Then, display the `.head()` of this new DataFrame.

```
# Your code here
grouped = df.groupby(['State', 'Gender'])['Deaths',
                                         'Population'].agg(['mean',
                                                             'min', 'max', 'std'])

grouped.head()
```



<style scoped> .dataframe tbody tr th:only-of-type { vertical-align: middle; }

```
.dataframe tbody tr th {
    vertical-align: top;
}

.dataframe thead tr th {
    text-align: left;
}

.dataframe thead tr:last-of-type th {
    text-align: right;
}
```

</style>

		Deaths				
		mean	min	max	std	mean
State	Gender					
Alabama	Female	10753.325000	10	116297	24612.250487	1.078713e
	Male	10765.850000	10	88930	20813.538537	1.014946e
Alaska	Female	679.975000	13	4727	1154.870455	1.440403e
	Male	860.357143	12	5185	1411.777392	1.518884e


```
# First, reset the index. Notice the subtle difference; State and Gender are now columns
grouped = grouped.reset_index()
grouped.head()
```

```
<style scoped> .dataframe tbody tr th:only-of-type { vertical-align: middle; }
```

```
.dataframe tbody tr th {
    vertical-align: top;
}
```

```
.dataframe thead tr th {
    text-align: left;
}
```

```
</style>
```

	State	Gender	Deaths				
			mean	min	max	std	
0	Alabama	Female	10753.325000	10	116297	24612.250487	1.078
1	Alabama	Male	10765.850000	10	88930	20813.538537	1.014
2	Alaska	Female	679.975000	13	4727	1154.870455	1.440
3	Alaska	Male	860.357143	12	5185	1411.777392	1.518
4	Arizona	Female	8998.386364	21	133923	26245.941003	1.246

Note how the way index is displayed has changed. The index columns that made up the multi-hierarchical index before are now stored as columns of data, with each row given a more traditional numerical index.

Let's confirm this by reexamining the `index` attribute of `grouped` in the cell below.

```
grouped.index
```

```
RangeIndex(start=0, stop=102, step=1)
```

However, look again at the displayed DataFrame -- specifically, the columns. Resetting the index has caused the DataFrame to use a multi-indexed structure for the columns.

In the cell below, examine the `columns` attribute of `grouped` to confirm this.

```
# Notice that this causes columns to be MultiIndexed!
grouped.columns

MultiIndex(levels=[['Deaths', 'Population', 'Gender', 'State'], ['mean', 'min',
'max', 'std', '' ]],
            codes=[[3, 2, 0, 0, 0, 0, 1, 1, 1, 1], [4, 4, 0, 1, 2, 3, 0, 1, 2,
3]])
```

Column levels

Since we're working with multi-hierarchical indices, we can examine the indices available at each level.

In the cell below, use the `.get_level_values()` method contained within the DataFrame's `columns` attribute to get the values for the outermost layer of the index.

```
grouped.columns.get_level_values(0)

Index(['State', 'Gender', 'Deaths', 'Deaths', 'Deaths', 'Deaths', 'Population',
      'Population', 'Population', 'Population'],
      dtype='object')
```

Now, get the level values for the inner layer of the index.

```
grouped.columns.get_level_values(1)

Index(['', '', 'mean', 'min', 'max', 'std', 'mean', 'min', 'max', 'std'],
      dtype='object')
```

Flattening the DataFrame

We can also **flatten** the DataFrame from a multi-hierarchical index to a more traditional one-dimensional index. We do this by creating each unique combination possible of every level of the multi-hierarchical index. Since this is a complex task, you do not need to write it -- but take some time to examine the code in the cell below and see if you can understand how it works!

```
# We could also flatten these:
cols0 = grouped.columns.get_level_values(0)
cols1 = grouped.columns.get_level_values(1)
grouped.columns = [col0 + '_' + col1 if col1 != '' else col0 for col0, col1 in list(
# The list comprehension above is more complicated then what we need but creates a n
# demonstrates using a conditional within a list comprehension.
# This simpler version works but has some tail underscores where col1 is blank:
# grouped.columns = [col0 + '_' + col1 for col0, col1 in list(zip(cols0, cols1))]
grouped.columns
```

```
Index(['State', 'Gender', 'Deaths_mean', 'Deaths_min', 'Deaths_max',
      'Deaths_std', 'Population_mean', 'Population_min', 'Population_max',
      'Population_std'],
      dtype='object')
```

Now that we've flattened the DataFrame, let's inspect a couple rows to see what it looks like.

In the cell below, inspect the `.head()` of the `grouped` DataFrame.

```
grouped.head()
```

```
<style scoped> .dataframe tbody tr th:only-of-type { vertical-align: middle; }
```

```
.dataframe tbody tr th {
    vertical-align: top;
}
```

```
.dataframe thead th {
    text-align: right;
}
```

```
</style>
```

	State	Gender	Deaths_mean	Deaths_min	Deaths_max	Deaths_

	State	Gender	Deaths_mean	Deaths_min	Deaths_max	Deaths_
0	Alabama	Female	10753.325000	10	116297	24612.250
1	Alabama	Male	10765.850000	10	88930	20813.538
2	Alaska	Female	679.975000	13	4727	1154.8704
3	Alaska	Male	860.357143	12	5185	1411.7773
4	Arizona	Female	8998.386364	21	133923	26245.947

Using pivots

Now, we'll gain some practice using the DataFrame's built-in `.pivot()` method.

In the cell below, call the DataFrame's `.pivot()` method with the following parameters:

- `index = 'State'`
- `columns = 'Gender'`
- `values = 'Deaths_mean'`

Then, display the `.head()` of our new `pivot` DataFrame to see what it looks like.

```
# Now it's time to pivot!
pivot = grouped.pivot(index='State', columns='Gender', values='Deaths_mean')
pivot.head()
```

```
<style scoped> .dataframe tbody tr th:only-of-type { vertical-align: middle; }

.dataframe tbody tr th {
    vertical-align: top;
}

.dataframe thead th {
    text-align: right;
}

</style>
```

Gender	Female	Male
State		

Gender	Female	Male
State		
Alabama	10753.325000	10765.850000
Alaska	679.975000	860.357143
Arizona	8998.386364	10036.204545
Arkansas	6621.615385	6301.690476
California	48312.840909	49555.522727

Great! We've just created a pivot table.

Let's reset the index and see how it changes our pivot table.

In the cell below, reset the index of the `pivot` object as we did previously. Then, display the `.head()` of the object to see if we can detect any changes.

```
# Again, notice the subtle difference of resetting the index:
pivot = pivot.reset_index( )
pivot.head()
```

```
<style scoped> .dataframe tbody tr th:only-of-type { vertical-align: middle; }

.dataframe tbody tr th {
    vertical-align: top;
}

.dataframe thead th {
    text-align: right;
}

</style>
```

Gender	State	Female	Male
0	Alabama	10753.325000	10765.850000
1	Alaska	679.975000	860.357143
2	Arizona	8998.386364	10036.204545
3	Arkansas	6621.615385	6301.690476

Gender	State	Female	Male
4	California	48312.840909	49555.522727

Visualizing Data With Pivot Tables


Now, we'll make use of our newly created pivot table to quickly create some visualizations of our data.

In the cell below, call `pivot.plot()` with the following parameters:

- `kind = 'barh'`
- `figsize = (15,8)`

```
# Now let's make a sweet bar chart!!
pivot.plot(kind='barh', figsize=(15,8))
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x10f6bd4a8>
```



Notice the Y-axis is currently just a list of numbers. That's because when we reset the index, it defaulted to assigning integers as the index for the DataFrame. Let's set the index back to 'State', and then recreate the visualization.


In the cell below:

- Use the `pivot` object's `.set_index()` method and set the index to `'State'`. Then, chain this with a `.plot()` call to recreate the visualization using the code we used in the cell above.

All the code in this cell should be done in a single line. Just call the methods -- do not rebind `pivot` to be equal to this line of code.

```
# Where's the states?! Notice the y-axis is just a list of numbers.
# This is populated by the DataFrame's index.
# When we used the .reset_index() method, we created a new numbered index to name ea
# Let's fix that by making state the index again.
pivot.set_index('State').plot(kind='barh', figsize=(15,8))
```

<matplotlib.axes._subplots.AxesSubplot at 0x1216805f8>



Now that we've created a visualization with the states as the y-axis, let's print out the head of the `pivot` object again.

```
# Also notice that if we call the DataFrame pivot again, state is not it's index.
# The above method returned a DataFrame with State as index and we plotted it,
# but it did not update the DataFrame itself.
pivot.head(2)
```

<style scoped> .dataframe tbody tr th:only-of-type { vertical-align: middle; }

```
.dataframe tbody tr th {  
    vertical-align: top;  
}
```

```
.dataframe thead th {  
    text-align: right;  
}
```

</style>

| Gender | State | Female | Male |
|--------|---------|-----------|--------------|
| 0 | Alabama | 10753.325 | 10765.850000 |
| 1 | Alaska | 679.975 | 860.357143 |

Note that the index has not changed. That's because the code we wrote when we set the index to the 'State' column returns a copy of the DataFrame object with the index set to 'State' -- by default, it does not mutate original `pivot` object.


If we want to do that, we'll need to capture the new object returned by updating the contents of the `pivot` variable.

In the cell below, set the index of `pivot` to 'State'. Then, recreate the bar plot using this new object.

```
# If we wanted to more permanently change the index we would set it first and then p  
pivot = pivot.set_index('State')  
pivot.plot(kind='barh', figsize=(15,8))
```



<matplotlib.axes._subplots.AxesSubplot at 0x121672f28>



Again, let's check the `.head()` of the DataFrame to confirm that the index structure has changed.

```
pivot.head(2)
```

```
<style scoped> .dataframe tbody tr th:only-of-type { vertical-align: middle; }

.dataframe tbody tr th {
    vertical-align: top;
}

.dataframe thead th {
    text-align: right;
}

</style>
```


| Gender | Female | Male |
|---------|-----------|--------------|
| State | | |
| Alabama | 10753.325 | 10765.850000 |
| Alaska | 679.975 | 860.357143 |

Finally, let's stack these bar charts to see how that looks.

In the cell below, recreate the visualization we did in the cell above, but this time, also pass in `stacked=True` as a parameter.

```
# Lastly, let's stack each of these bars for each state.
# Notice we don't have to worry about index here, because we've already set it above
pivot.plot(kind='barh', figsize=(15,8), stacked=True)
```

<matplotlib.axes._subplots.AxesSubplot at 0x121ccf4a8>



Stacking and Unstacking DataFrames

Now, let's get some practice stacking and unstacking DataFrames.

Stacking

In the cell below, let's display the head of `grouped` to remind ourselves of the format we left it in.

```
grouped.head()
```

```
<style scoped> .dataframe tbody tr th:only-of-type { vertical-align: middle; }
```

```
.dataframe tbody tr th {
    vertical-align: top;
}
```

```
.dataframe thead th {
```

```
text-align: right;
}
```

```
</style>
```

| | State | Gender | Deaths_mean | Deaths_min | Deaths_max | Deaths_std |
|---|---------|--------|--------------|------------|------------|--------------|
| 0 | Alabama | Female | 10753.325000 | 10 | 116297 | 24612.250000 |
| 1 | Alabama | Male | 10765.850000 | 10 | 88930 | 20813.530000 |
| 2 | Alaska | Female | 679.975000 | 13 | 4727 | 1154.870000 |
| 3 | Alaska | Male | 860.357143 | 12 | 5185 | 1411.777000 |
| 4 | Arizona | Female | 8998.386364 | 21 | 133923 | 26245.940000 |

As we can see above, `grouped` is currently in a flattened format, with no hierarchical structure to it's indices.

In the cell below, call the `grouped` `DataFrame`'s `.stack()` method.

```
grouped.stack()
```

| | | |
|---|-----------------|-------------|
| 0 | State | Alabama |
| | Gender | Female |
| | Deaths_mean | 10753.3 |
| | Deaths_min | 10 |
| | Deaths_max | 116297 |
| | Deaths_std | 24612.3 |
| | Population_mean | 1.07871e+06 |
| | Population_min | 2087 |
| | Population_max | 4334752 |
| | Population_std | 1.40031e+06 |
| 1 | State | Alabama |
| | Gender | Male |
| | Deaths_mean | 10765.9 |
| | Deaths_min | 10 |
| | Deaths_max | 88930 |
| | Deaths_std | 20813.5 |
| | Population_mean | 1.01495e+06 |
| | Population_min | 1129 |
| | Population_max | 4284775 |
| | Population_std | 1.39783e+06 |
| 2 | State | Alaska |
| | Gender | Female |

```

Deaths_mean      679.975
Deaths_min        13
Deaths_max       4727
Deaths_std       1154.87
Population_mean   144040
Population_min    1224
Population_max    682855
Population_std    201579
...
99  State      Wisconsin
    Gender      Male
    Deaths_mean  9573.45
    Deaths_min   13
    Deaths_max   113692
    Deaths_std   25681.4
    Population_mean  1.13532e+06
    Population_min  1286
    Population_max  6860107
    Population_std  2.08907e+06
100 State      Wyoming
    Gender      Female
    Deaths_mean  1161.03
    Deaths_min   10
    Deaths_max   13140
    Deaths_std   2937.94
    Population_mean  146757
    Population_min  336
    Population_max  672620
    Population_std  235238
101 State      Wyoming
    Gender      Male
    Deaths_mean  1149.51
    Deaths_min   10
    Deaths_max   10113
    Deaths_std   2569.28
    Population_mean  139224
    Population_min  244
    Population_max  694760
    Population_std  241360

```

Length: 1020, dtype: object

As we can see, the `.stack()` method has stacked our DataFrame from a flattened format into one with a multi-hierarchical index! This is an easy, quick way to aggregate our data.

Unstacking

Now, we'll explore unstacking with the `pivot` DataFrame, which is already stacked into a pivot table.

In the cell below, set `unstack` `pivot` using the object's `.unstack()` method. Then, display the object to see how it has changed.

```
pivot = pivot.unstack()
pivot
```

| Gender | State | |
|--------|----------------------|--------------|
| Female | Alabama | 10753.325000 |
| | Alaska | 679.975000 |
| | Arizona | 8998.386364 |
| | Arkansas | 6621.615385 |
| | California | 48312.840909 |
| | Colorado | 6460.162791 |
| | Connecticut | 7144.641026 |
| | Delaware | 2000.029412 |
| | District of Columbia | 1497.580645 |
| | Florida | 36019.071429 |
| | Georgia | 15372.317073 |
| | Hawaii | 2182.944444 |
| | Idaho | 2874.323529 |
| | Illinois | 23432.926829 |
| | Indiana | 13425.717949 |
| | Iowa | 6419.707317 |
| | Kansas | 5492.309524 |
| | Kentucky | 10426.083333 |
| | Louisiana | 9076.585366 |
| | Maine | 3471.823529 |
| | Maryland | 9894.780488 |
| | Massachusetts | 13356.846154 |
| | Michigan | 18421.659091 |
| | Minnesota | 8168.204545 |
| | Mississippi | 6342.634146 |
| | Missouri | 12493.170732 |
| | Montana | 2341.393939 |
| | Nebraska | 3667.794872 |
| | Nevada | 3729.166667 |
| | New Hampshire | 3293.344828 |
| | ... | |
| Male | Massachusetts | 11368.341463 |
| | Michigan | 17940.431818 |
| | Minnesota | 7792.795455 |
| | Mississippi | 6487.317073 |
| | Missouri | 11810.119048 |

| | |
|----------------|--------------|
| Montana | 2081.102564 |
| Nebraska | 3290.682927 |
| Nevada | 4489.261905 |
| New Hampshire | 2800.303030 |
| New Jersey | 15085.317073 |
| New Mexico | 3549.428571 |
| New York | 29864.477273 |
| North Carolina | 15750.409091 |
| North Dakota | 1587.411765 |
| Ohio | 23551.951220 |
| Oklahoma | 7468.909091 |
| Oregon | 6528.977273 |
| Pennsylvania | 27187.463415 |
| Rhode Island | 2239.243243 |
| South Carolina | 9078.292683 |
| South Dakota | 1800.500000 |
| Tennessee | 13333.050000 |
| Texas | 34347.636364 |
| Utah | 3081.511628 |
| Vermont | 1785.846154 |
| Virginia | 12585.833333 |
| Washington | 9877.431818 |
| West Virginia | 6211.612903 |
| Wisconsin | 9573.454545 |
| Wyoming | 1149.514286 |

Length: 102, dtype: float64

Note that it has unstacked the multi-hierarchical structure of the `pivot` DataFrame by one level. Let's call it one more time and display the results!

In the cell below, set `pivot` equal to `pivot.unstack()` again, and then print the `pivot` object to see how things have changed.

```
pivot = pivot.unstack()
pivot
```

```
<style scoped> .dataframe tbody tr th:only-of-type { vertical-align: middle; }
```

```
.dataframe tbody tr th {
    vertical-align: top;
}
```

```
.dataframe thead th {
    text-align: right;
}
```

</style>

State	Alabama	Alaska	Arizona	Arkansas	California
Gender					
Female	10753.325	679.975000	8998.386364	6621.615385	48312.840909
Male	10765.850	860.357143	10036.204545	6301.690476	49555.522727

2 rows × 51 columns

After calling `unstack` a second time, we can see that `pivot` has a flattened structure since it has been completely unstacked!

Summary

In this lab, we learned how to:

- Use `.groupby()` to stack and slice data conditionally
- Use aggregate methods in combination with groupby statements
- Create pivot tables with pandas
- Leverage pivot tables and groupby statements to create quick visualizations
- `stack` and `unstack` DataFrames

Releases

No releases published

Packages

No packages published

Contributors 5



Languages

● Jupyter Notebook 100.0%